

SHIYI CAO

shiyicao314@gmail.com ♦ +86 18190733275

EDUCATION

ETH Zurich

Sept. 2020 - May. 2023

Department of Computer Science

M.S. in Computer Science

Relevant Courses: Advanced System Lab (6/6), Cloud Computing Architecture (5.75/6), Information Security Lab (5.75/6), Design of Parallel and High-Performance Computing (5.5/6), Advanced Operating System (5.25/6)

Shanghai Jiao Tong University

Sept. 2016 - 2020

School of Electronic Information and Electrical Engineering

B.S. in Computer Science and Technology

GPA: 88/100

Relevant Courses: Computer Architecture Lab (94), Computing Theory (97), Mathematical Foundations of Computer Science (92), Operating Systems Lab (90), Cloud Computing (93)

RESEARCH INTEREST

My general research interests lie in the fields of computer systems and high-performance computing, with a focus on understanding and accelerating emerging applications on heterogeneous systems.

SELECTED PUBLICATIONS

1. **Shiyi Cao**, Salvatore Di Girolamo, and Torsten Hoefer. Accelerating Data Serialization/Deserialization Protocols with In-Network Compute. In *Workshop on Exascale MPI, ExaMPI@SC*, 2022.
2. **Shiyi Cao**, Yuanning Gao, Xiaofeng Gao, and Guihai Chen. Adam: An adaptive fine-grained scheme for distributed metadata management. In *International Conference on Parallel Processing (ICPP)*, 2019.

RESEARCH EXPERIENCE

Generalized Pipeline Parallelism for DL Model Training

Current

Catalyst, CMU, Advisor: Zhihao Jia

- Leading the end-to-end implementation for enabling generalized parallelism training strategies.

Accelerating Data SerDes with In-Network Compute

Aug. 2021 - Aug. 2022

SPCL Lab, ETH Zurich, Advisor: Salvatore Di Girolamo and Prof. Torsten Hoefer

- Offloaded data deserialization (ProtoBuf) to SmartNIC for efficient RPC framework.
- Designed the deserialization algorithm to enable parallel and streaming processing on the NIC.
- Published the paper on SC'22 ExaMPI Workshop as the *first author*.

Deep Reinforcement Learning in Distributed Metadata Management

Sept. 2018 - Jan. 2019

Advanced Network Laboratory, SJTU, Advisor: Prof. Xiaofeng Gao

- Introduced for the first time deep reinforcement learning in distributed metadata management.
- Proposed an adaptive fine-grained metadata management scheme AdaM, leveraging deep reinforcement learning.
- Conducted experiments on real-world data traces and compared AdaM with strong baselines.

- Demonstrated that AdaM can address the trade-off between load balance and locality preservation cost-effectively and is highly adaptive to time-varying access pattern.
- Published the paper in ICPP'19 (International Conference on Parallel Processing) as the *first author*.

SELECTED PROJECTS

Barrelfish OS Development

Mar. 2022 - Jun. 2022

Advanced Operating System Course by David Cock and Prof. Timothy Roscoe

- Implemented our own memory management, paging, message passing, inter-core communication etc. on Barrelfish research operating system.
- Implemented and benchmarked the Network stack.

Distributed DL Training on Bagua

Oct. 2021 - Jan. 2022

DS3Lab, ETH, Advisor: Jiawei Jiang and Prof. Ce Zhang

- Port, improve and benchmark existing distributed deep compression training algorithms to Bagua, a deep learning training acceleration framework for PyTorch.

High-performance Image Compression Implementation

Mar. 2021 - June. 2021

Advanced System Lab Course Project

- Designed highly optimized implementations of the whole SPIHT image compression pipeline, leveraging techniques such as SIMD vectorization, memory rearrangement, and blocking.
- Our best optimized version achieves a runtime speedup of 100x and 200x for encoding and decoding respectively compared with the baseline implementation.

TALKS & PRESENTATIONS

Participant, Workshop on Exascale MPI @ SC

Nov. 2022

- Made oral presentation for the accepted paper Accelerating Data Serialization/Deserialization Protocols with In-Network Compute.

Participant, International Conference on Parallel Processing

Aug. 2019

- Made oral presentation for the accepted paper Adam: An adaptive fine-grained scheme for distributed metadata management.

AWARDS

- Academic Excellence Scholarship (Second Class), 2016-2017
- Academic Excellence Scholarship (Third Class), 2018-2019
- Meng Minwei International Exchange Fund (12000RMB), 2019

SKILLS

English Proficiency

GRE: 329 + 4.0 (V:160 Q:169 AW:4.0), TOEFL: 109

Programming

C, C++, SSE/AVX, Python, PyTorch, Tensorflow

Softwares

Latex, Matlab, Unity3D

GitHub

<https://github.com/caoshiyi>